

# Get the best out of Live Sessions HOW?

# e!



## Check your Internet Connection

Log in 10 mins before, and check your internet connection to avoid any network issues during the LIVE session.

## Speak with the Instructor

By default, you will be on mute to avoid any background noise. However, if required you will be **unmuted by instructor**.



## Clear Your Doubts

Feel free to clear your doubts. Use the “Questions” tab on your webinar tool to interact with the instructor at any point during the class.

## Let us know if you liked our content

Please share feedback after each class. It will help us to enhance your learning experience.



edureka!



# DevOps Certification Training

# COURSE OUTLINE

## MODULE 09

1. Introduction to DevOps

2. Version Control with Git

3. Git and Jenkins

4. Continuous Integration with Jenkins

5. Configuration Management using Ansible

6. Containerization using Docker Part - I



7. Containerization using Docker Part - II

8. Container Orchestration Using  
Kubernetes Part-I

9. Container Orchestration Using  
Kubernetes Part-II

10. Monitoring Using Prometheus and  
Grafana

11. Provisioning Infrastructure using  
Terraform Part - I

12. Provisioning Infrastructure using  
Terraform Part - II

**edureka!**

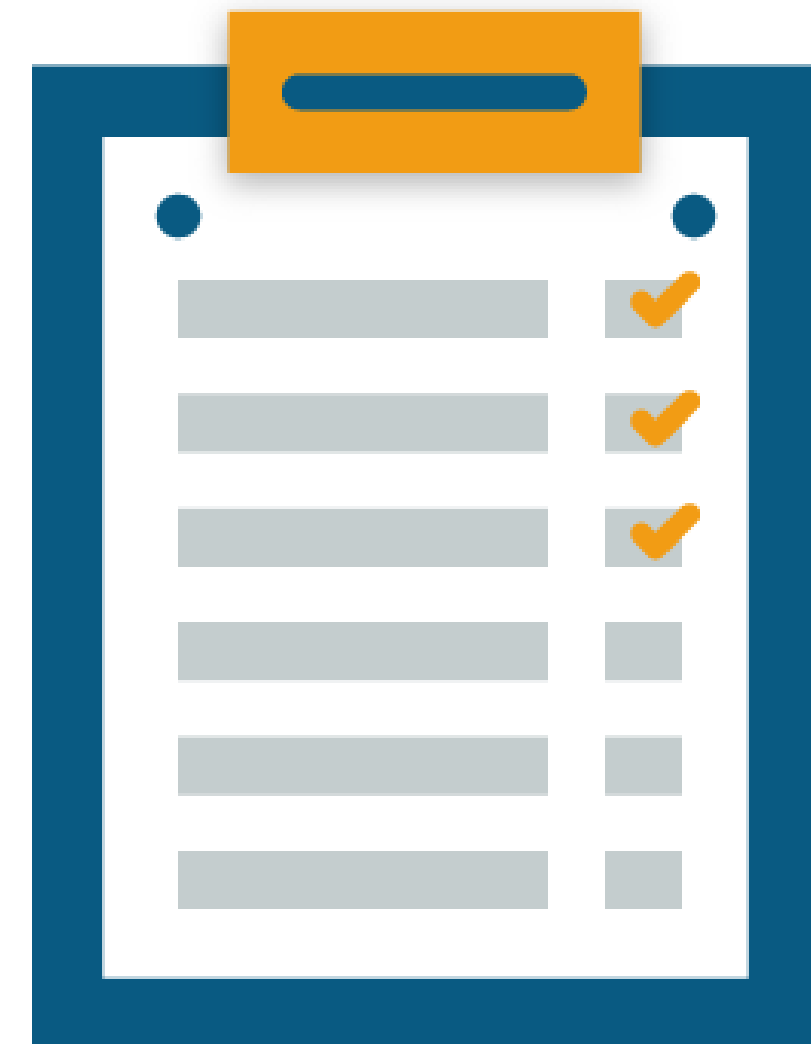
# **Module 9 – Container Orchestration Using Kubernetes Part-II**

# Topics

---

Following are the topics covered in this module:

- Services
- Persistent Storage in Kubernetes
- Primitives for PersistentVolumeClaims
- Secrets and ConfigMaps
- Headless Services
- StatefulSets
- Helm Charts





# Objectives

---

After completing this module, you should be able to:

- Deploy different Kubernetes Services
- Utilize Volumes to store Persistent Data
- Create Persistent Volume Claims to attach volumes to Pods
- Understand Persistent Volume Claims Primitives
- Use Headless Services in Stateful Sets
- Deploy Helm Charts





# Labels, Selectors and Annotations

# Labels


---

Key-value pairs attached to Kubernetes objects such as nodes, pods, and controllers

Used to organise objects by identifying similar attributes

Added at the time of object creation

Can be added/modified later in the object's lifecycle





# Selectors


---

Core grouping mechanism in Kubernetes

Users/applications can pick a group of objects by their labels using selectors

Multiple labels can be used with selectors using a comma separator, which acts as a logical AND (&&) operator

Different APIs have context-based implementation for empty selectors



# Types of Selectors

---

Currently, there are two types of selectors supported by Kubernetes API:



# Equality-based Selector

---

- Allows filtering by key-value matching
- All the specified labels must be satisfied by the matching object
- Two kinds of operators are used:
  - `== ( = )` : represents equality
  - `!=` : represents inequality



# Equality-based Selector (Contd.)

- Eg : env = production
  - This selects objects having “production” as the value for the key “env”
  - env != production inversely selects all the objects not having “production” as their value for “env” key

## From our example:

```
$ kubectl get pods --selector owner=edureka
```

NAME	READY	STATUS	RESTARTS	AGE
pi-144jr	0/1	Completed	0	1h



# Set-based Selectors

---

- Allows filtering the keys based on a set of values
- Supports three kinds of operators
  - in
  - notin
  - exists
- Eg: team in (HR, training )
  - Finds objects having “HR” or “training” values for the key “team”



# Annotations

---

Attaches non-identifying metadata to objects

The metadata can be utilised by various tools and libraries

Uses key-value pair with valid segments i.e. an optional prefix and name separated by / (slash)

Metadata can have characters not permitted by labels





# Annotations (Contd.)

---

Annotations can store information such as:

- Build/ release/ image information
- Logging and monitoring pointers
- Analytics or audit repositories
- Library/ tool information
- Information for external ecosystems
- Custom Declarative Configuration layer fields




# Services

# Services

---


- Services load-balance the requests in and out of the application
- It acts as a layer of abstraction between the containers and the network



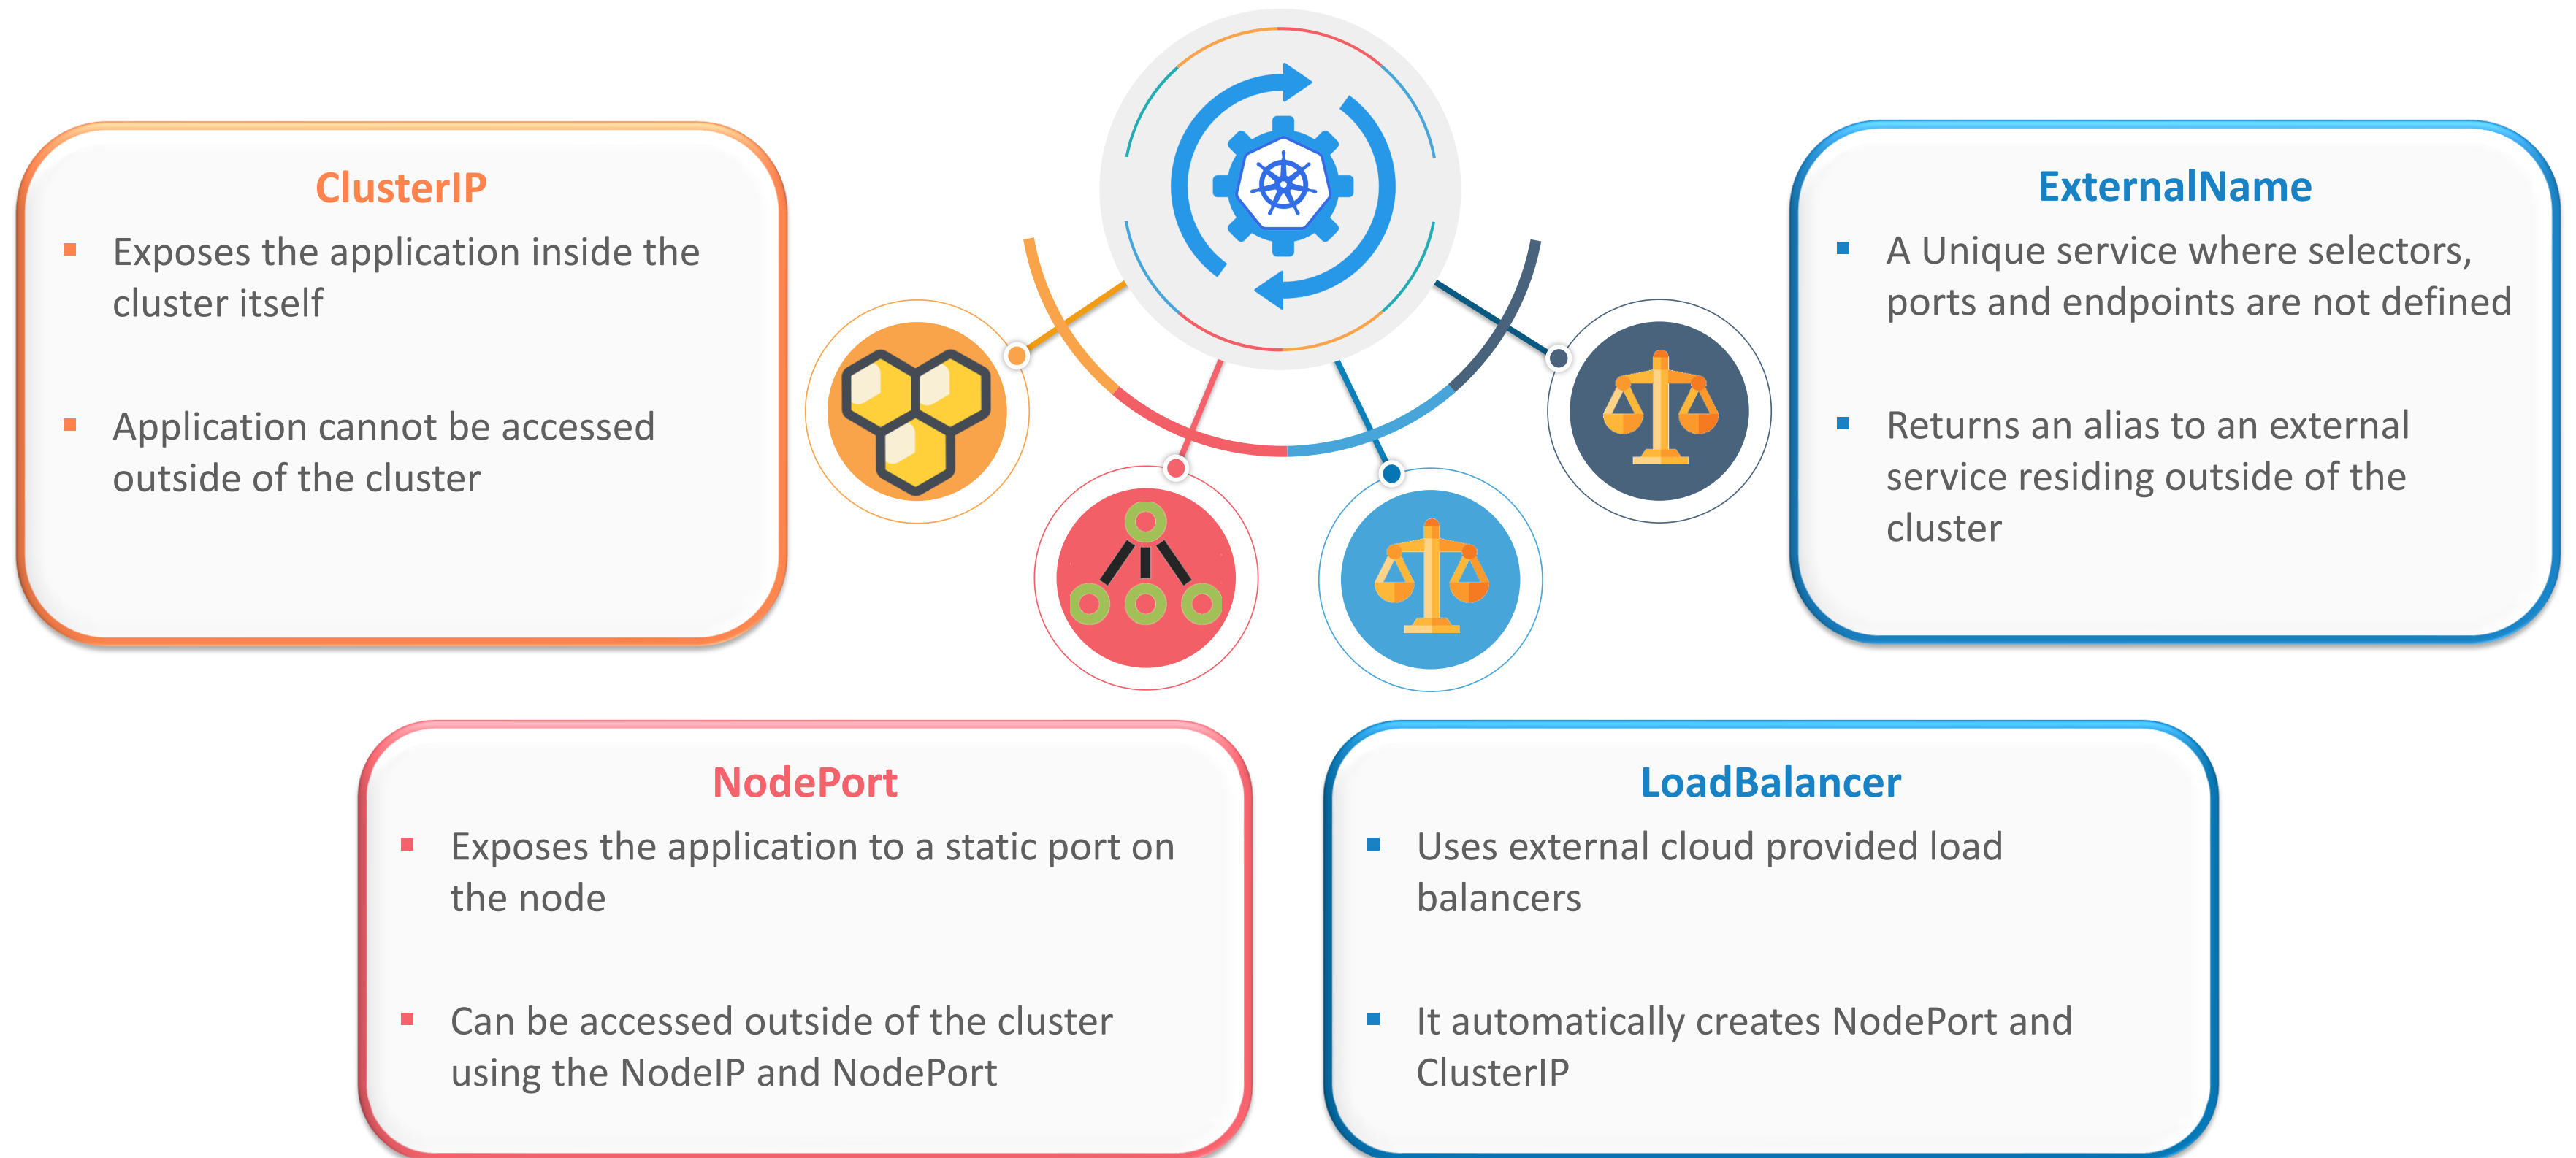
# Services

---

- 01 Services act as a single constant point of entry to a group of pods lending similar service
- 02 It solves pod related issues such as pods being ephemeral and horizontal scaling
- 03 Each service has an immutable IP address and port that will never change throughout the service's lifecycle
- 04 Clients directly communicate with the service. This allows pods to be scheduled anywhere
- 05 It gives resurrecting pods more consistency than a deployment




# Types of Services



# Types of Services (Contd.)

---



- Services are superset of each other
- NodePort is superset of ClusterIP and Loadbalancer is superset of NodePort




# Demo: Services

# Persistent Volumes

# Volumes: The Problem

---


- 01 Normal volumes provide persistent storage but the layer of abstraction is lost
- 02 Difficult to manage tightly coupled applications which share volume storage
- 03 Volumes are bound to pods lifecycle and are destroyed once the pod is terminated
- 04 Unable to provide persistent storage to stateful/ distributed application across the cluster



# Persistent Storage in Kubernetes

---


- 01 The layer of abstraction can be bought back using Persistent Volumes and Persistent Volume Claims
- 02 Persistent Volumes can be utilised as a cluster resource
- 03 Kubernetes binds the Persistent Volume Claim to the best possible Persistent Volume
- 04 The Volume Claim connects the applications with the persistent storage
- 05 Persistent Volumes provide the stability required for stateful and distributed applications



# Persistent Storage in Kubernetes

---

- Kubernetes has the provision of storing Persistent Data in Volumes
- This maintains consistency for stateful applications



# Persistent Volumes and Persistent Volume Claims

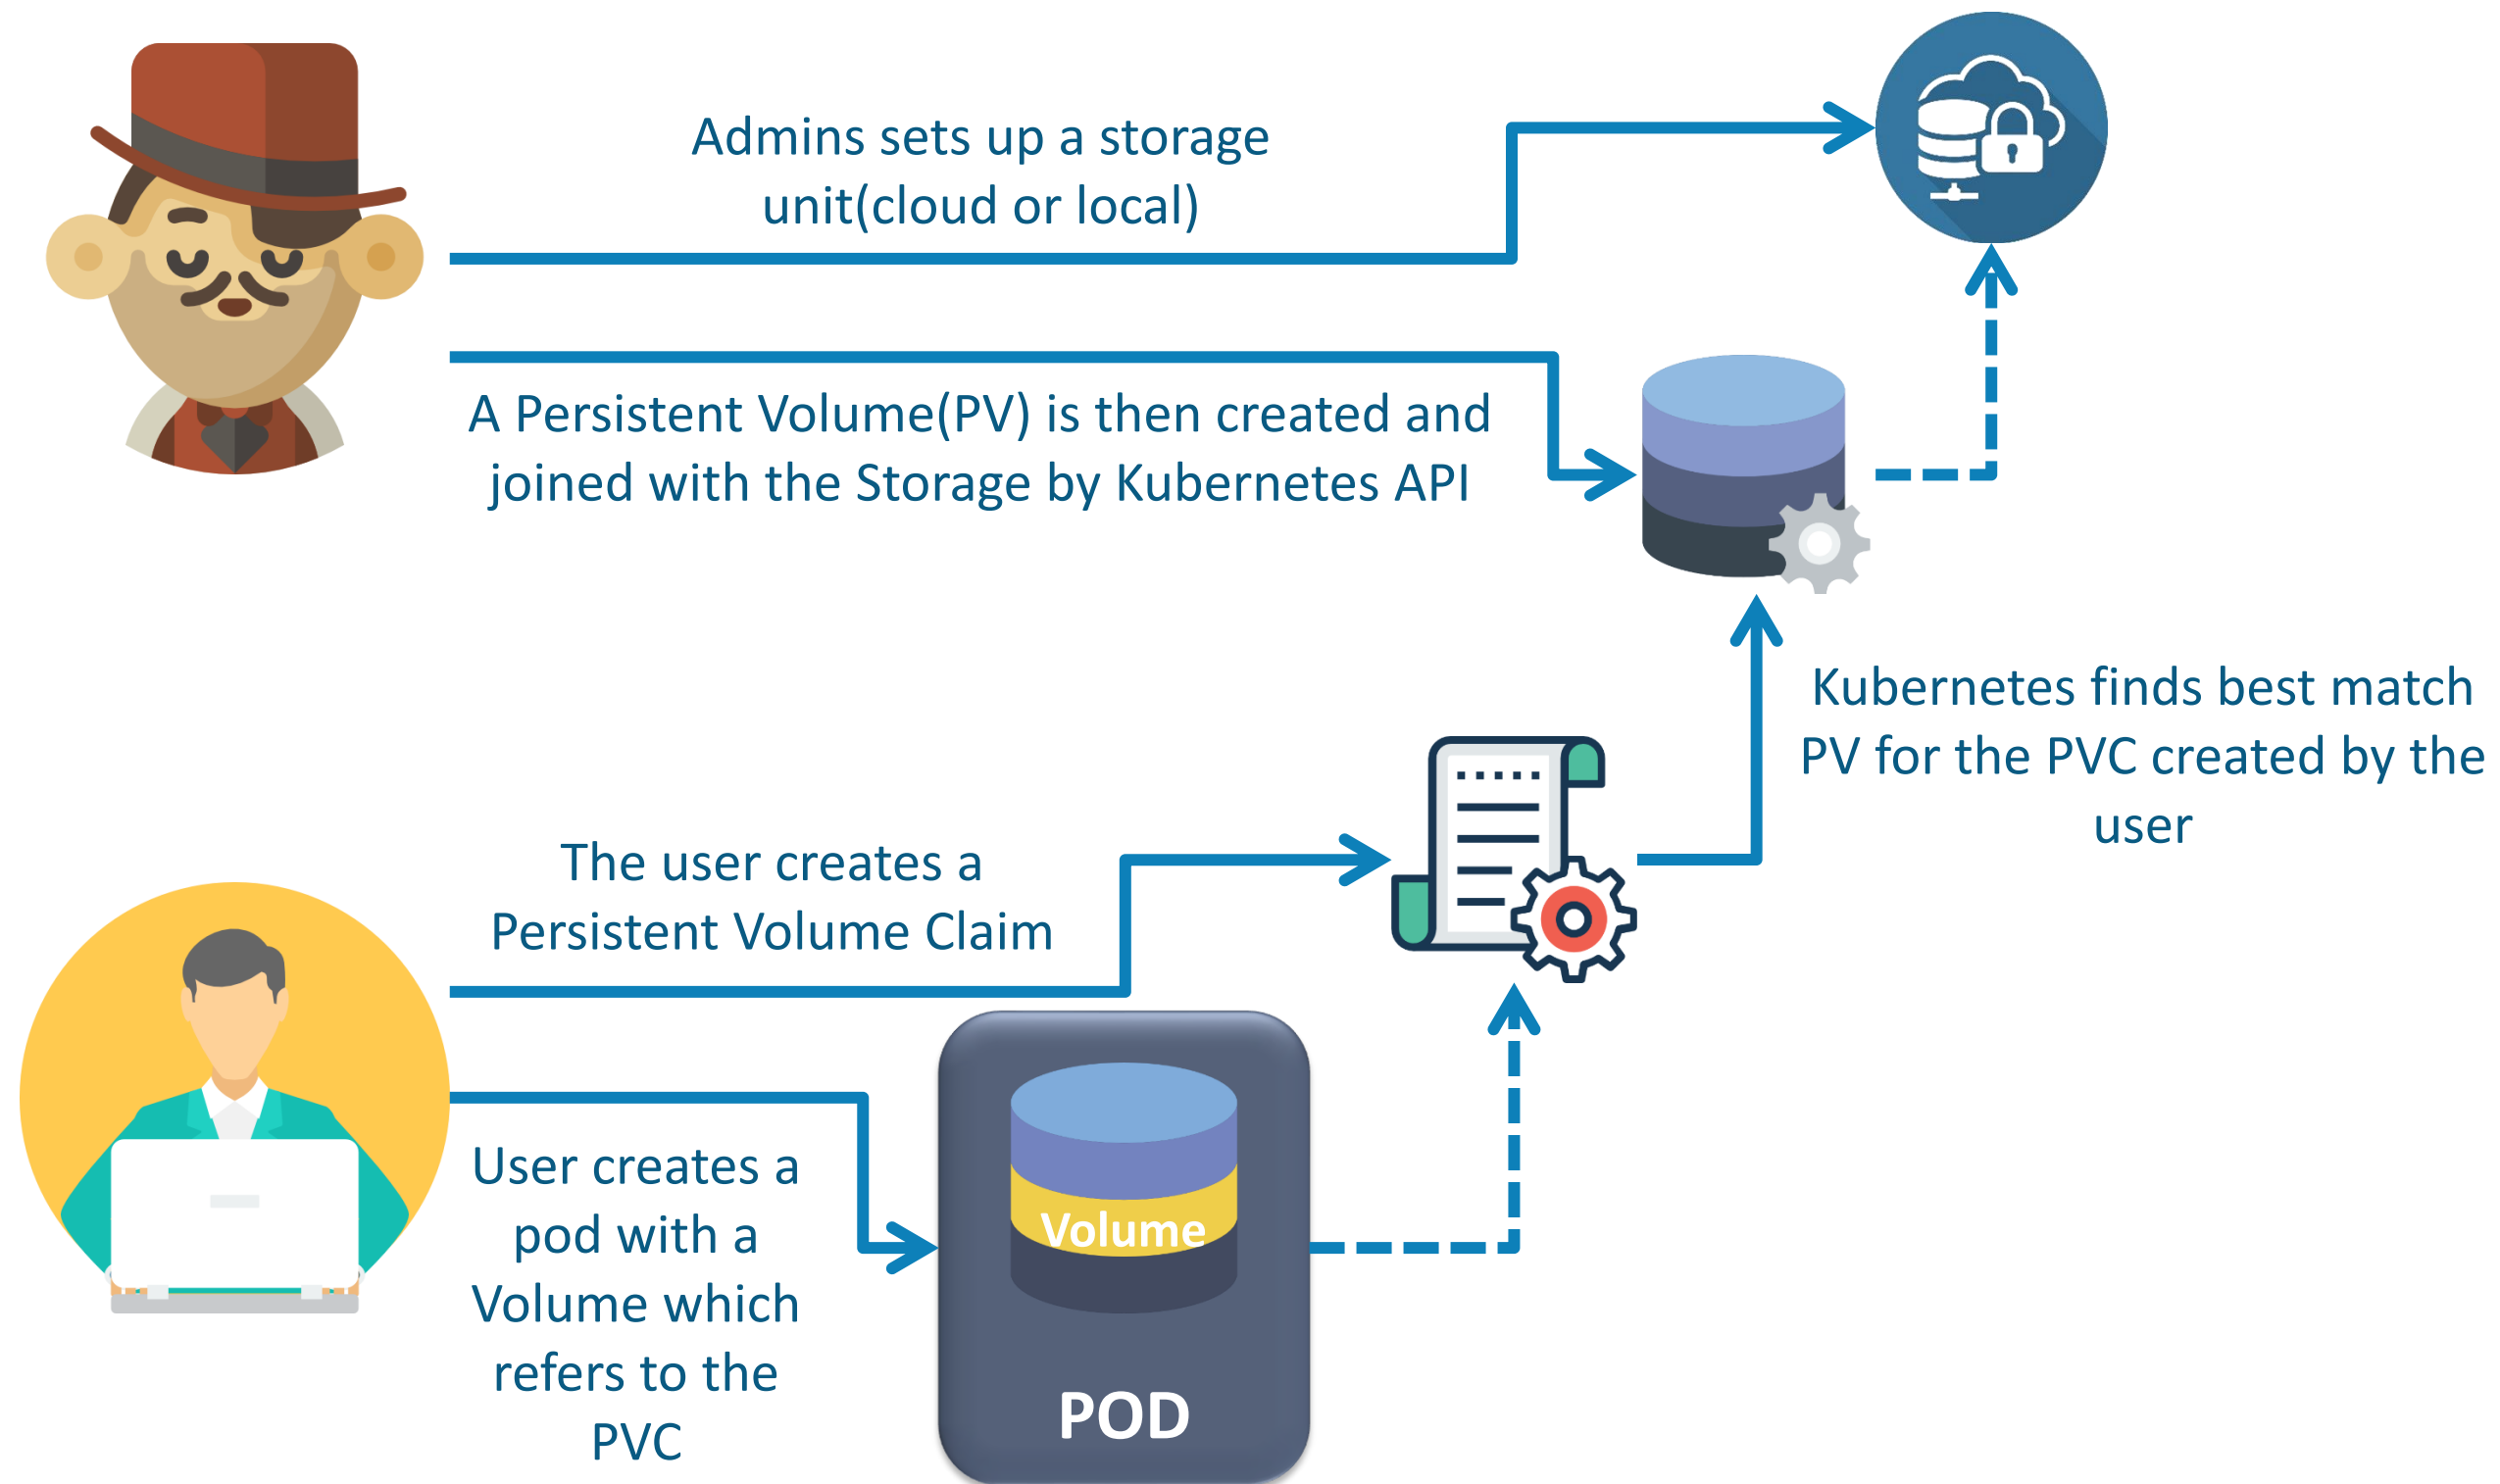
---

- 01 Persistent Volume once created, acts as a resource in the cluster
- 02 Persistent Volume Claims are requests for the created resources(PVs)
- 03 PVCs ensure that the application will be able to use the required storage
- 04 PVs and PVCs abstract the storage implementation from the application
- 05 The application needs to specify certain parameters for its storage requirement to the PVC





# Persistent Volumes and Persistent Volume Claims





# Access Modes for Persistent Volumes

# Access Modes for Persistent Volumes

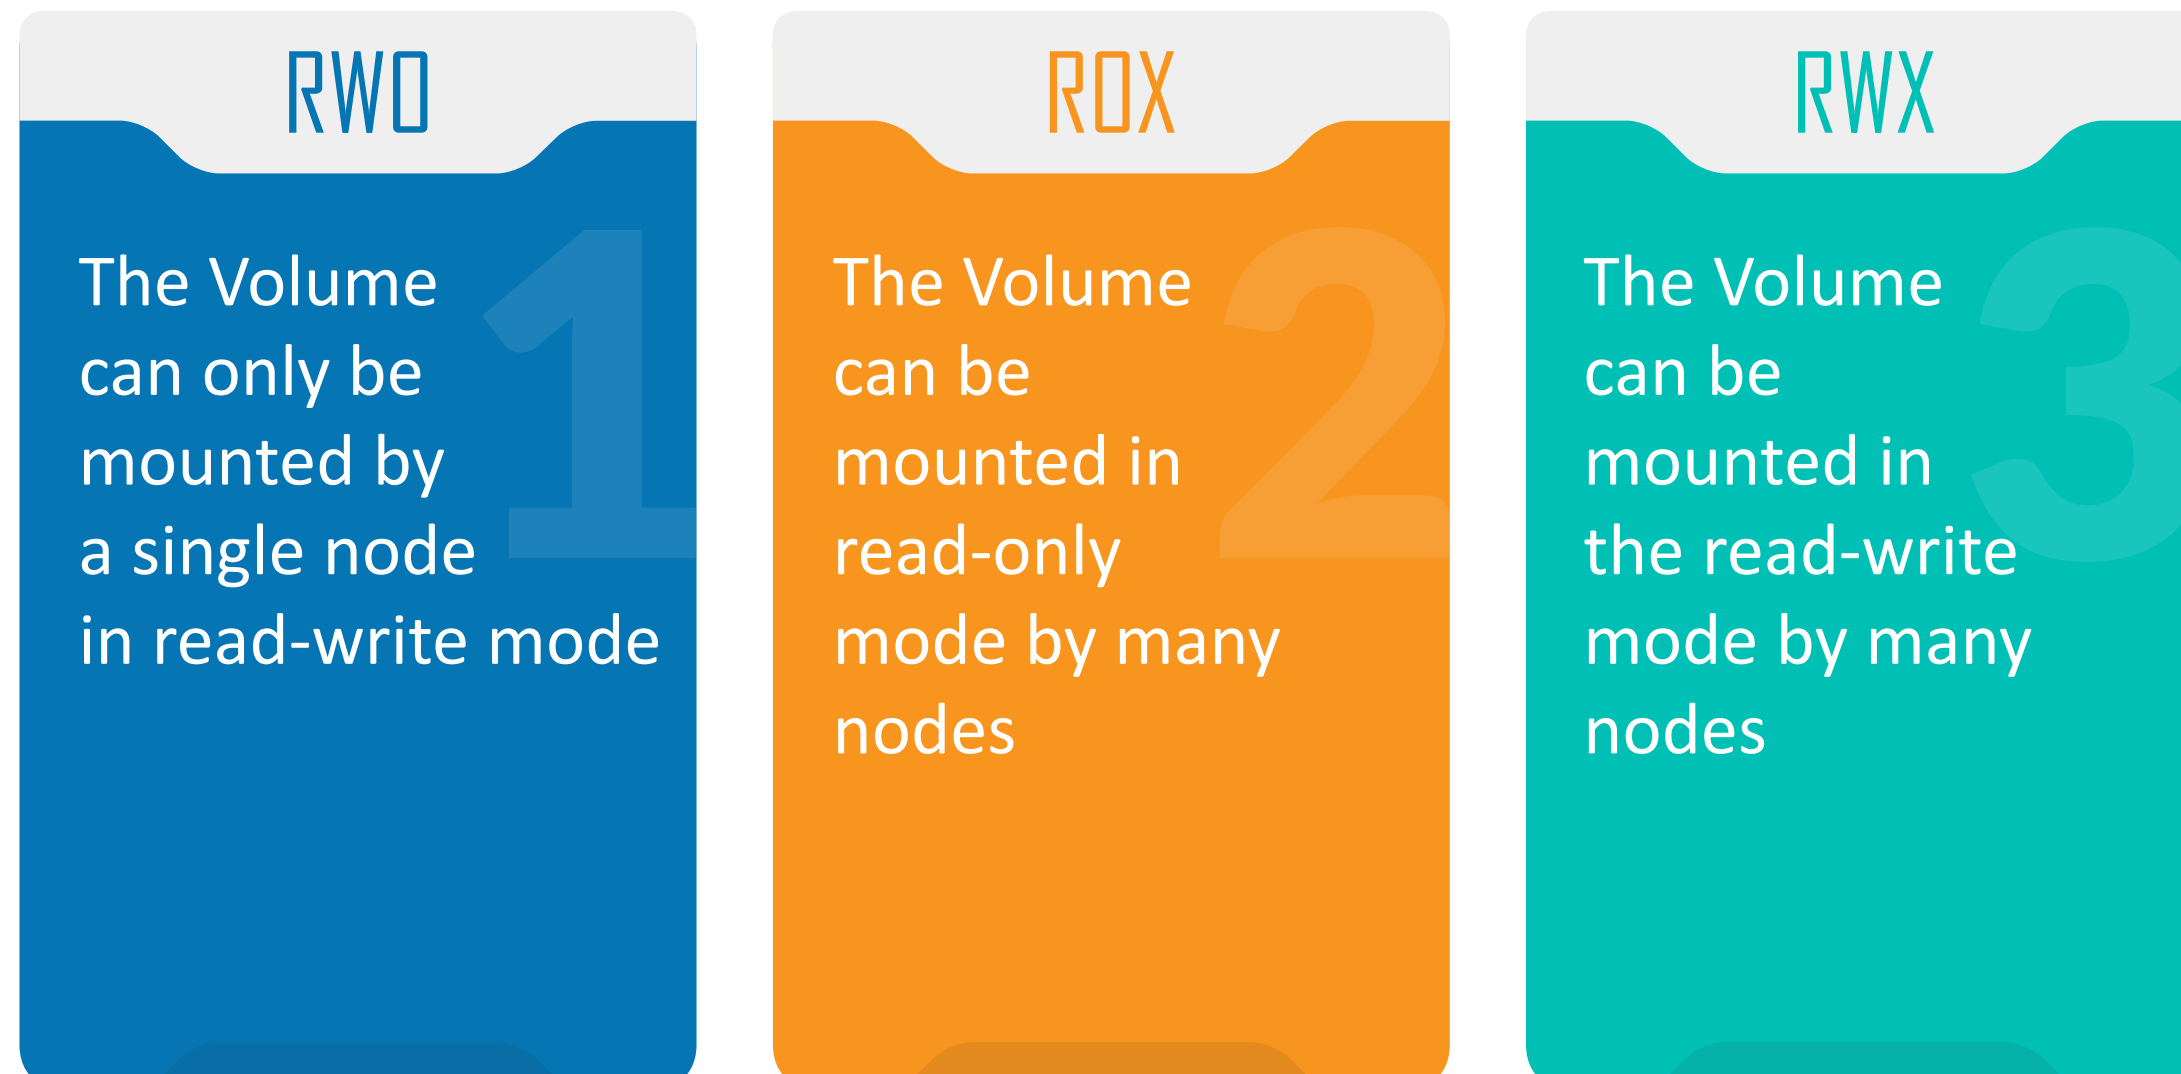
---

There are three types of Access Modes for Persistent Volumes:



# Access Modes for Persistent Volumes (Contd.)

---

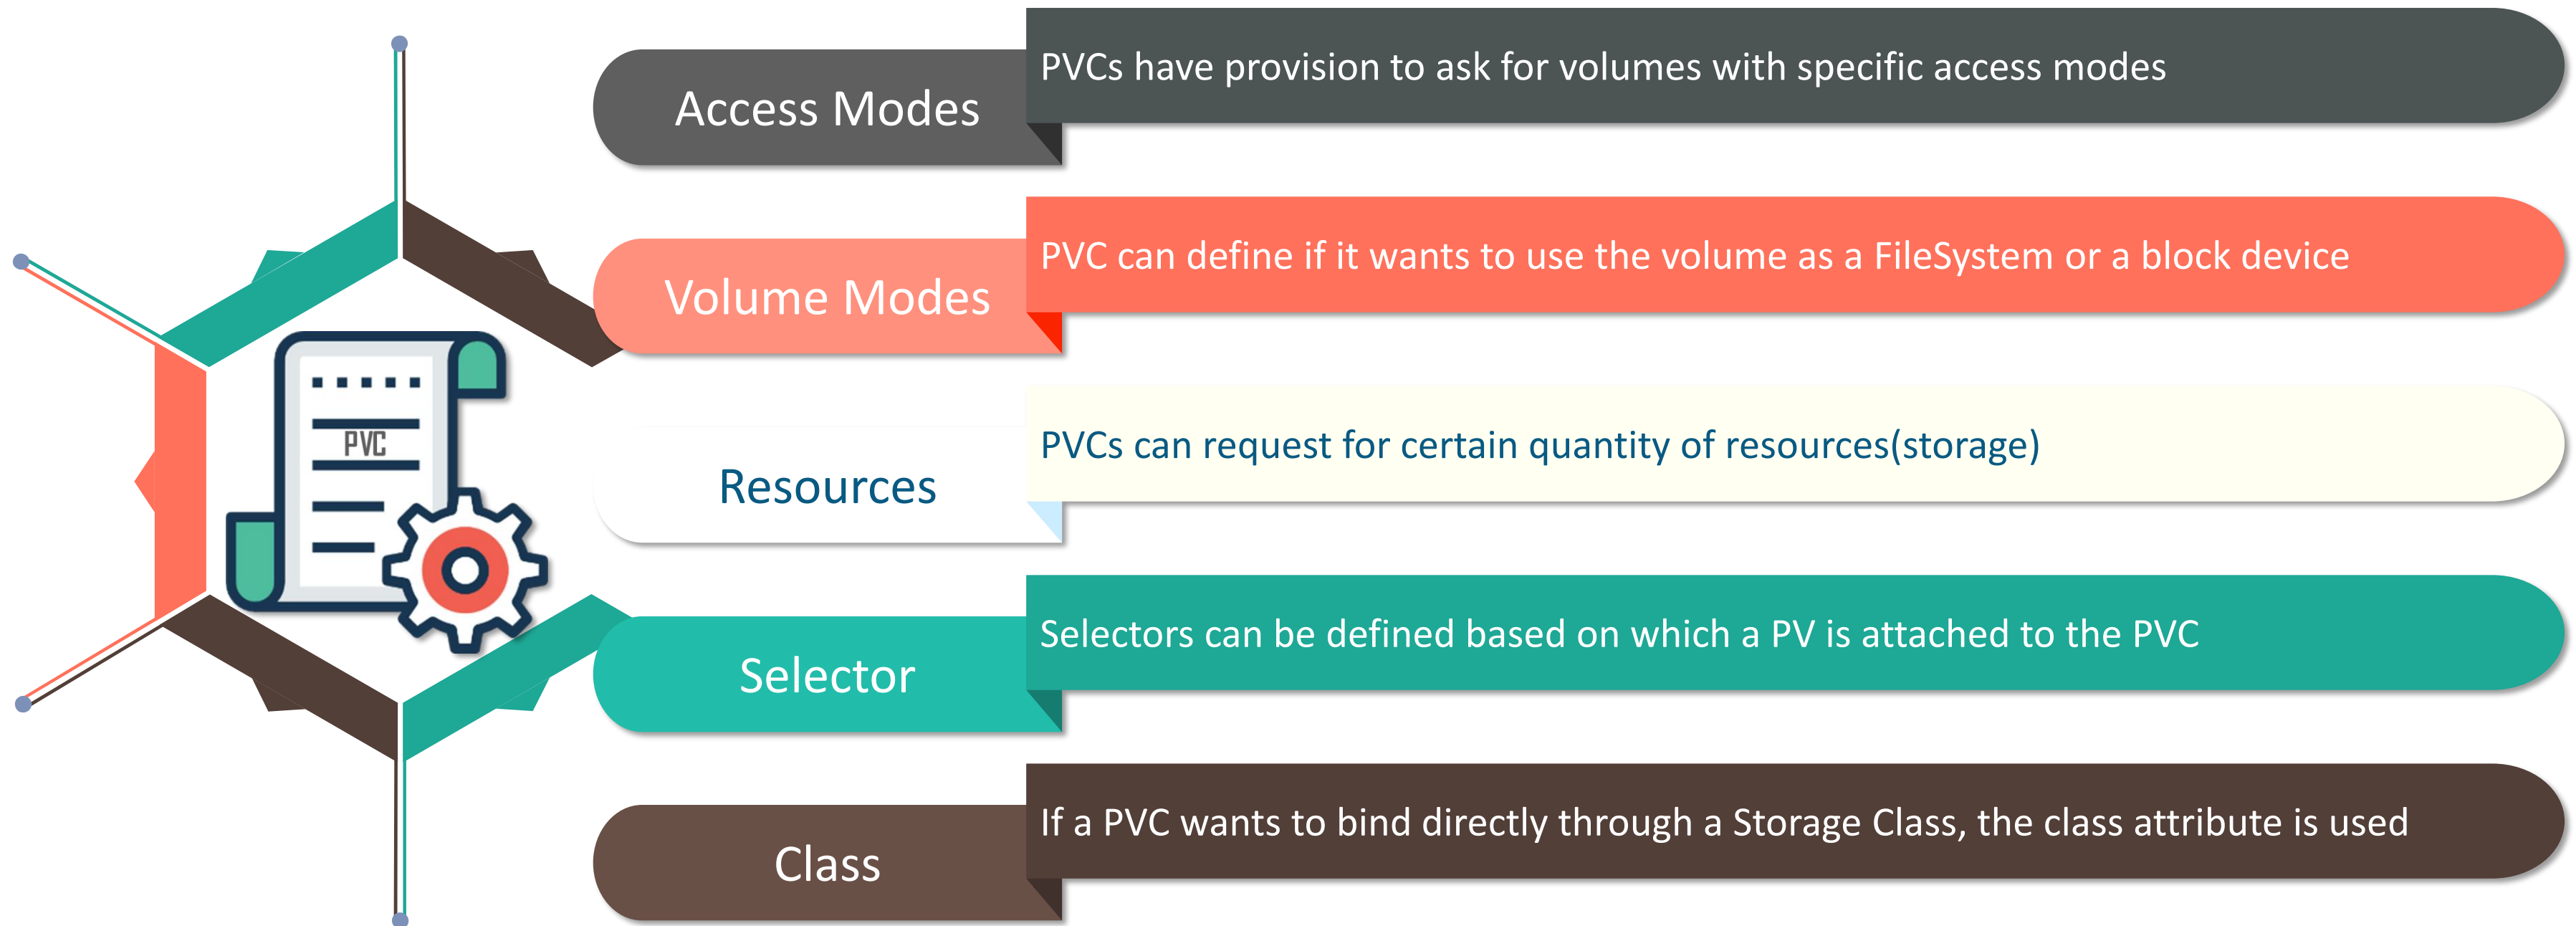




# Persistent Volume Claims Primitives

# PVC Primitives

---








# Demo: Persistent Volumes and Persistent Volume Claims

# Headless Service

# Headless Service


---

A service with its ClusterIP set to **none** is known as a Headless Service



**Note:** Unlike other services, headless service does not provide any load-balancing or proxy capabilities

# Headless Service (Contd.)



Helps decouple service from Kubernetes giving an alternate form of service discovery

DNS entries are still recorded on the cluster to discover pods through the service

Headless Service helps the client to connect to all the pods

StatefulSet requires a headless service for its pods to be discovered by the network


Running a DNS query for a Headless Service returns the list of pods attached to it

# StatefulSets

# StatefulSets

---

- StatefulSets create pods with sticky identities
- The pods retain their identity through various reschedulings



# StatefulSets (Contd.)


01 StatefulSets provides a stable identity and state for non-fungible applications

02 Keeps track of the number of pods deployed and ensures that the said number is maintained

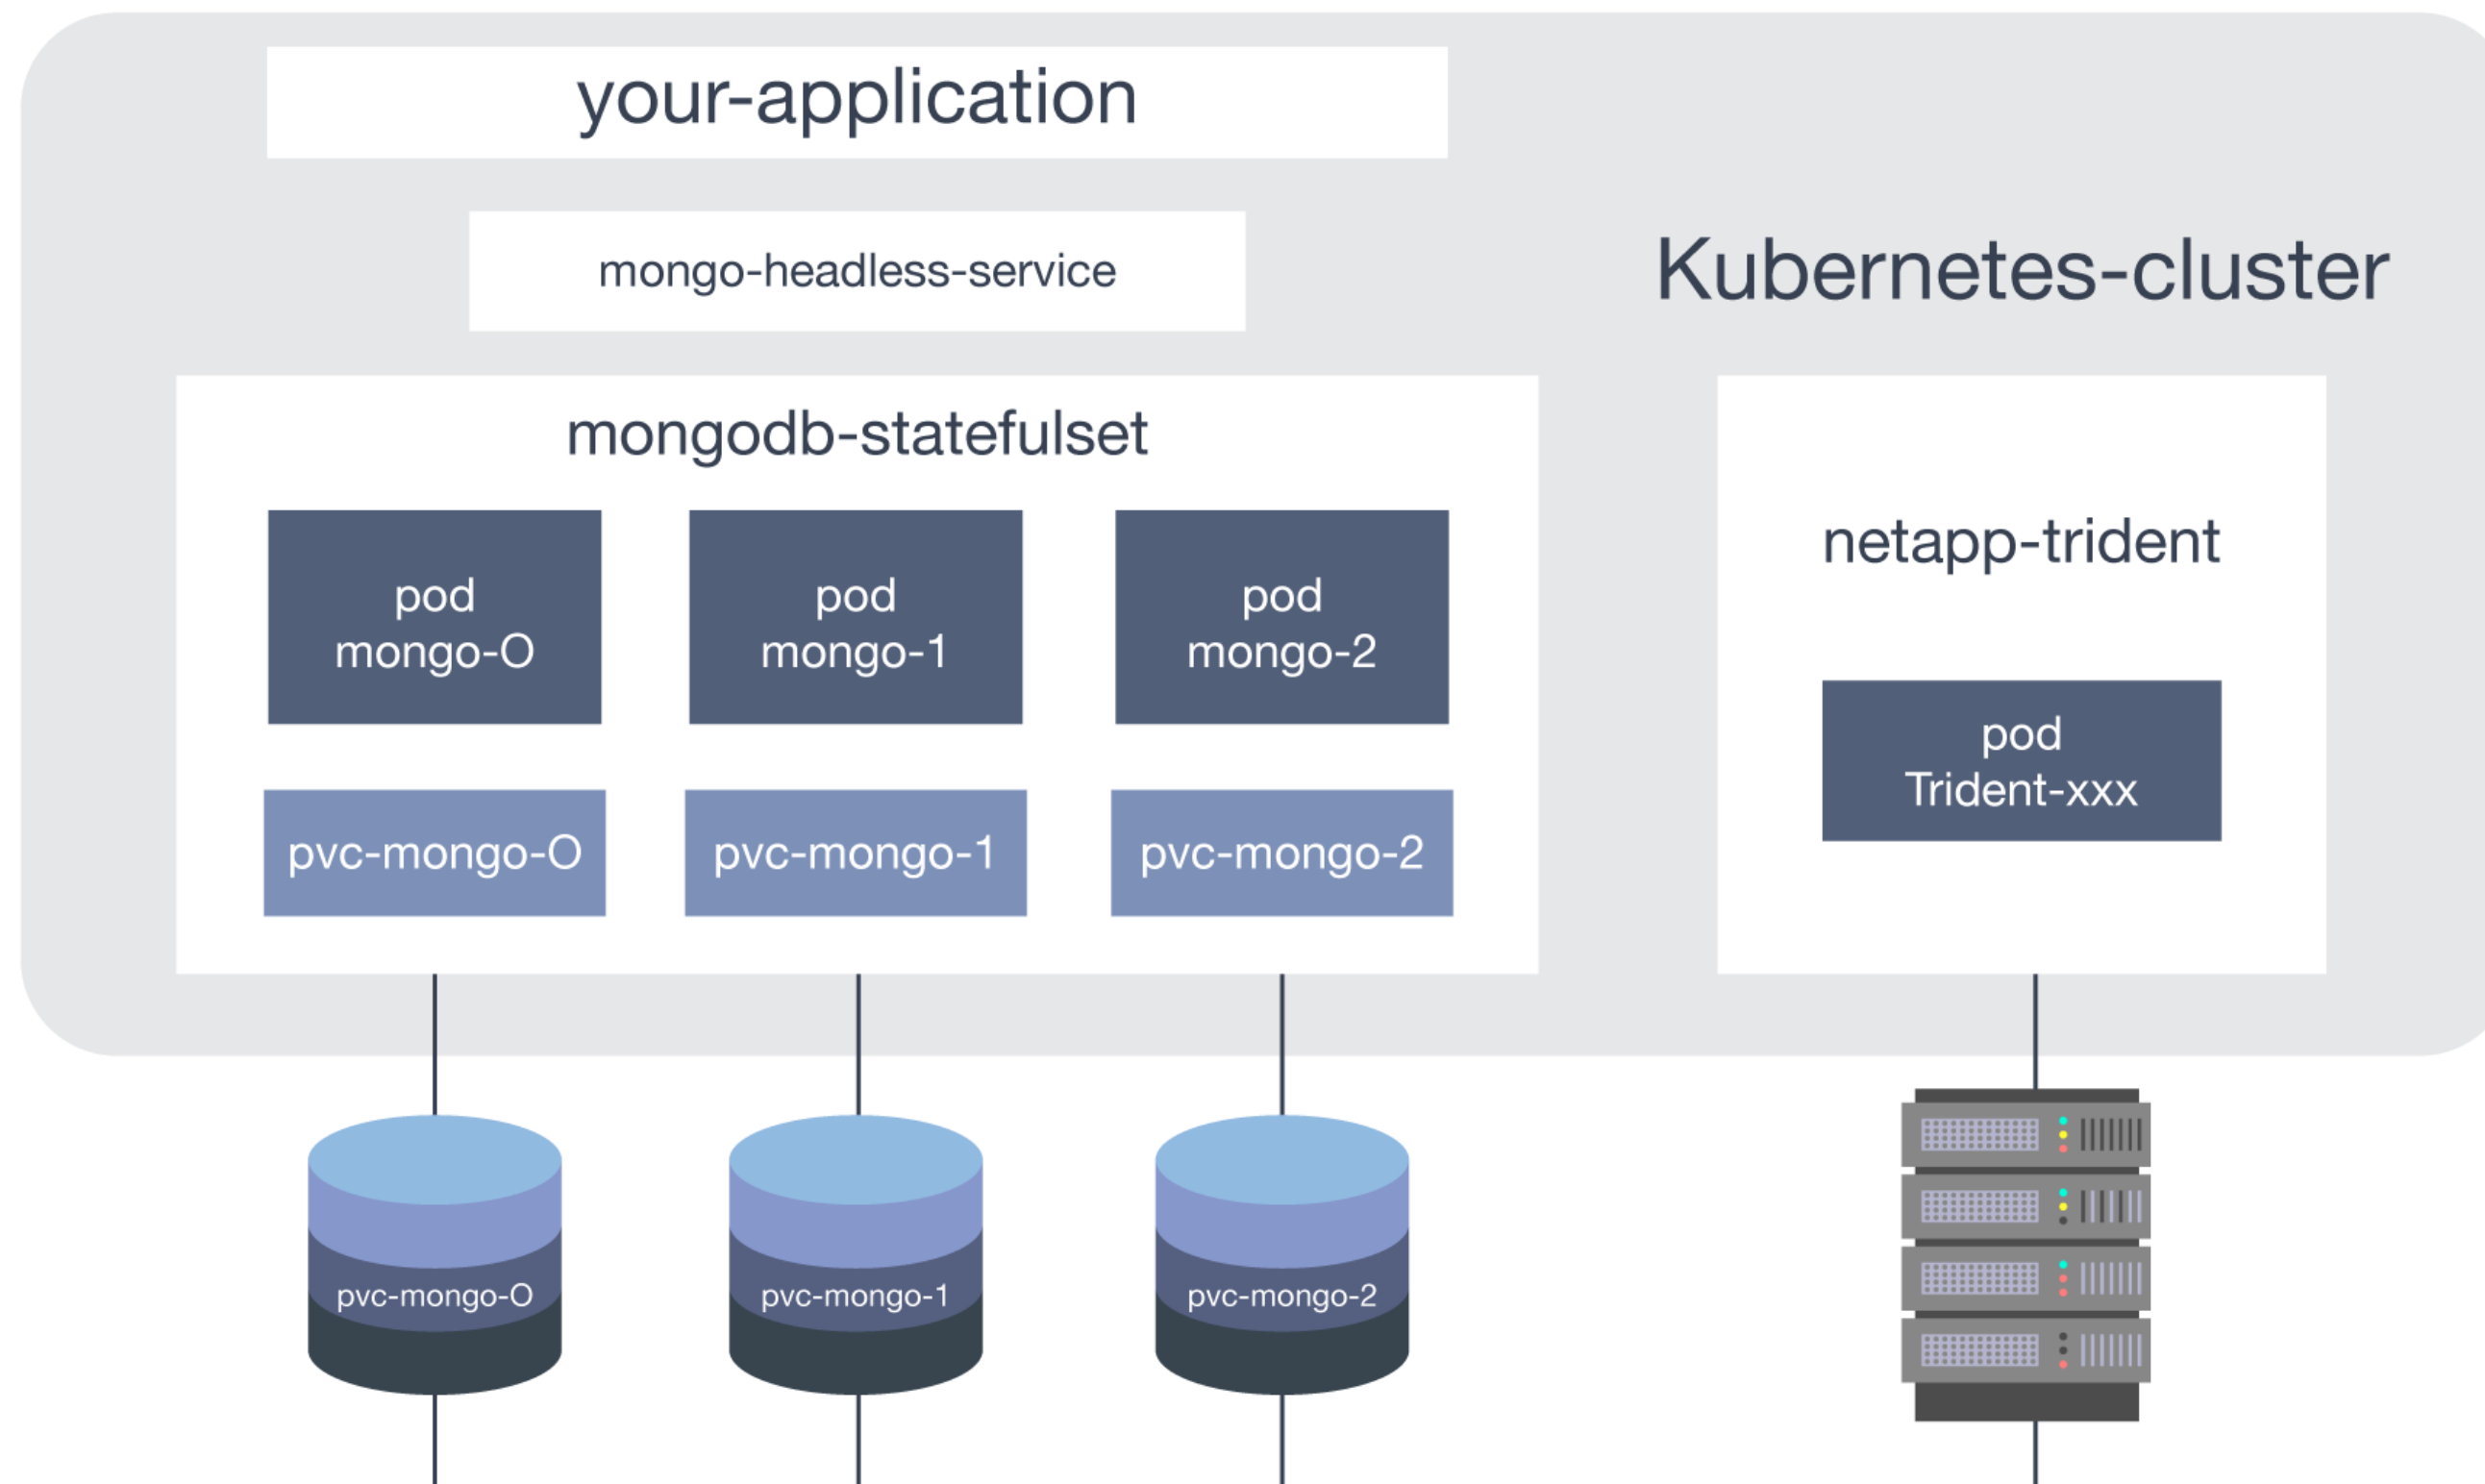
03 Pod template is specified as a part of StatefulSet's specification

04 The Pods scheduled are not exact replicas of each other. Each pod has its own storage unit to work with

05 Rolling-updates can be applied to pods associated with StatefulSets



# Sample Use Case for Stateful Applications







# Demo: StatefulSets

# ConfigMaps and Secrets

# ConfigMaps

---


**01** Exports configuration data into a container to avoid hardcoding it inside the application

**02** Employs the ConfigMap API resource which contains configuration data in the form of key-value pairs

**03** Used to pass non-sensitive information to the containers

**04** The contents of the ConfigMap are not directly passed to the application


**05** Content is passed either to the container as environment variables or to the volumes as files



# Secrets

---

- 01 Secrets allow passing sensitive information to the containers
- 02 Holds information in key-value pairs and can be passed as environment variables or files in a volume
- 03 Only distributed to the nodes that are running the pod which need access to the secrets content
- 04 Stored on memory (RAM) of the node instead of the physical memory (HDD or SSD)
- 05 On the master node, secrets are stored in encrypted form on the etcd storage





# Demo: ConfigMaps and Secrets



# Kubernetes Package Manager: Helm

# Helm

---

- Packages collection of Kubernetes logic which can be deployed later
- The packages can be distributed and reused as HELM Charts



# Helm (Contd.)

---



Official Cloud Native Foundation supported open-source Kubernetes project

It can install software, its dependencies, update existing software, configure deployments and so on

Uses a command line utility called **helm** and a companion component called **tiller**

**Tiller** runs on the cluster to handle configurations and deployments. It constantly listens for incoming **helm** commands

Stores packages in the form of helm charts and maintains a chart repository

**Note:** Tiller has been removed from Helm Version3



# Helm Charts

---

- Consists of YAML configurations and templates
- Configurations are passed as manifest files to Kubernetes

Directory structure of a Helm Chart:

- package-name/
  - charts/
  - templates/
  - Chart.yaml
  - LICENCE
  - README.md
  - requirements.yaml
  - values.yaml



# Helm Charts (Contd.)



**charts/**

Dependencies which are managed manually are placed in this directory

**templates/**

Contains config files combined with values which are tuned into Kubernetes manifests

**Chart.yaml**

Holds metadata about the chart with information such as chart name, version, keywords and so on

**LICENCE**

Plaintext file containing the licencing information of the helm chart's content

**README.md**

Readme information document for the helm chart

**requirements.yaml**

Consists of the chart's dependencies

**Values.yaml**

Holds the default configuration values for the chart

# Demo: Helm Charts

# Summary

## Services

- Services load-balance the requests in and out of the application
- It acts as a layer of abstraction between the containers and the network




edureka!

Copyright © edureka and/or its affiliates. All rights reserved.

## Persistent Storage in Kubernetes

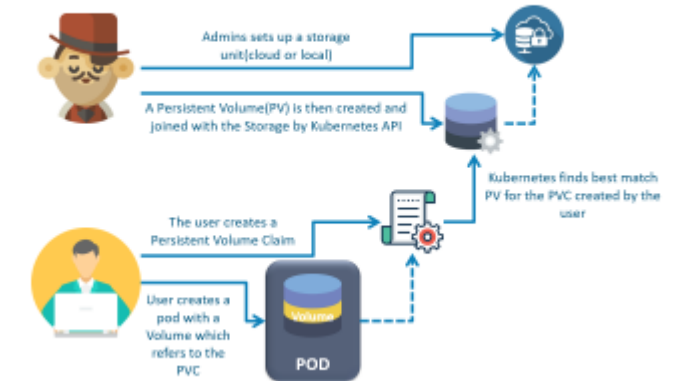
- Kubernetes has the provision of storing Persistent Data in Volumes
- This maintains consistency for stateful applications



edureka!

Copyright © edureka and/or its affiliates. All rights reserved.

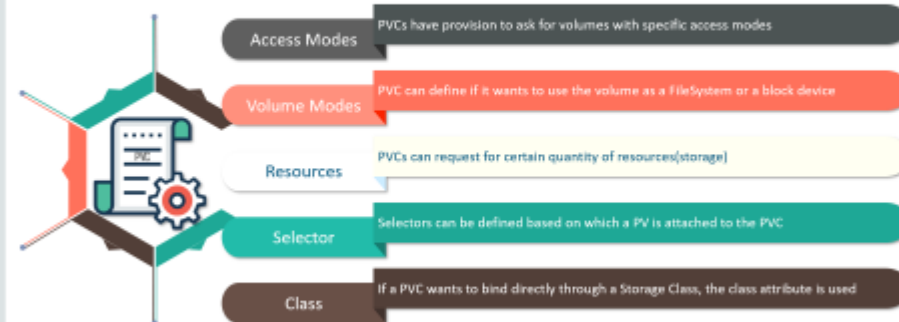
## Persistent Volumes and Persistent Volume Claims



edureka!

Copyright © edureka and/or its affiliates. All rights reserved.

## PVC Primitives



edureka!

Copyright © edureka and/or its affiliates. All rights reserved.

## StatefulSets

- StatefulSets create pods with sticky identities
- The pods retain their identity through various reschedulings



edureka!

Copyright © edureka and/or its affiliates. All rights reserved.

## Helm

- Packages collection of Kubernetes logic which can be deployed later
- The packages can be distributed and reused as HELM Charts



edureka!

Copyright © edureka and/or its affiliates. All rights reserved.

# Questions



# FEEDBACK







# Thank You

---

For more information please visit our website  
[www.edureka.co](http://www.edureka.co)